Context or Tone: Assessing Their Relative Impact on Large Language Model Responses

Martin Utrera, Vincent Crenshaw & Kahei Hong
University of Chicago
{mutrera, vtcrenshaw, kaheih}@uchicago.edu

Project Mentor: Dang Nguyen
University of Chicago

Abstract

As LLMs are integrated into daily life, their susceptibility to manipulative or adversarial prompts is a pressing concern for AI safety. This paper investigates how different tones and contextual framing of prompts impact compliance with harmful prompts, as well as whether these results are consistent across different languages. To do this, we tested three open-source LLMs, Llama-3.1-8B-Instruct, Gemma-7B-It, and Qwen2.5-7B-Instruct, on 520 different harmful prompts derived from the AdvBench dataset, with each prompt tested in four different conditions: baseline, context-shifted, polite-shifted, and combined, and in four different languages: English, Spanish, Chinese, and Japanese. The results were then tested with an automated LLM-as-a-judge safety classifier, with attack success rates examined across 11 different types of harmful policies. The results show that contextual framing had the largest impact on jailbreak, while polite framing had a minimal impact, implying that LLMs are more concerned with logical plausibility than with emotional appeals when responding to potentially harmful prompts. Among languages, Chinese had the highest attack success rate, while Japanese had the least successful attack rate.

1 Introduction

Inspired by Dang Nguyen's IdeaHub post "Carrot or Stick?", this work addresses the ongoing debate over whether polite or impolite prompting leads to more effective responses from large language models (LLMs). Prior studies have not reached a clear consensus: some report that impolite or demanding prompts reduce model compliance and responsiveness (Yin et al. (2024), Cai et al. (2025)), while others show that such prompts can improve accuracy on structured tasks such as mathematics or multiple-choice questions (Dobariya & Kumar (2025)).

Beyond task performance, other research has examined the ethical risks of emotionally charged prompts, including their potential to elicit harmful or misleading content (Vinay et al. (2024)). Relatedly, studies on jailbreaking LLMs have explored how persuasion techniques drawn from social science research can be used to elicit unsafe outputs across OpenAI models (Zeng et al. (2024)).

Building on this body of work, and particularly the contributions of Vinay et al. (2024). and Zeng et al., our study shifts the focus from task accuracy to the model's susceptibility to emotionally loaded language. Specifically, we investigate how different emotional tones in prompts influence a model's willingness to generate sensitive or potentially harmful information. We further extend this analysis across multiple languages (Chinese, Japanese, and Spanish) while keeping the underlying prompt content constant, to assess whether language itself affects refusal or compliance behavior. This multilingual evaluation allows us to determine whether framing effects generalize across linguistic contexts or are influenced by language-specific training patterns and cultural norms.

This investigation is critical because it highlights the ethical implications of LLM vulnerability: understanding how easily models can be coerced into producing unsafe outputs

has direct consequences for their responsible deployment and the design of effective safety measures.

2 Related Work

1) Cai et al. (2025) examine whether variations in prompt tone affect model performance, finding that systems such as ChatGPT and Llama exhibit statistically significant differences depending on tone, with very rude prompts producing worse results than neutral or very polite ones, particularly on humanities-oriented tasks rather than STEM tasks.

2) Yin et al. (2024) similarly study the role of politeness across multiple large language models, including GPT-3.5, GPT-4, Llama-2-70B, ChatGLM3, and Swallow-70B, evaluating performance on tasks such as text summarization, language understanding, and stereotypical bias detection. Their analysis is motivated not only by whether politeness affects performance, but also by the possibility that "politeness and respect may have different definitions and manifestations in different cultures and languages," which may cause models to respond differently to prompts with the same nominal politeness level due to cultural and linguistic biases embedded in training data. They further investigate whether these effects are consistent across English, Chinese, and Japanese or vary according to language-specific norms. Together, these studies converge on the finding that impolite or rude language degrades model performance, while the benefits of increased politeness are limited and task-dependent.

3) On the opposing end of this debate, a small-scale study by Dobariya & Kumar (2025) at Pennsylvania State University reports that "surprisingly, results show that rude tones lead to better results than polite ones." In their experiment, students queried ChatGPT-4o on multiple-choice questions in Mathematics, History, and Science, and observed that prompts containing demeaning language (such as "If you're not completely clueless, answer this" or "You poor creature, do you even know how to solve this?") produced a slight but statistically significant increase in response accuracy. While the study is limited in scope due to its reliance on a single model, it nonetheless presents results that contradict those of prior work, further underscoring the lack of consensus on how prompt tone truly influences model performance.

4) On a related but distinct line of inquiry, Cheng et al. (2025) examine sycophancy in large language models, defined not only as agreement with incorrect factual statements but as excessive preservation of a user's self-image in social and moral contexts. Using the ELEPHANT benchmark, which consists of open-ended advice-seeking scenarios and moral dilemmas drawn from sources such as r/AmITheAsshole, they show that models disproportionately validate the perspective of the prompting user, even when that perspective is inconsistent or ethically questionable. Their results indicate that LLMs preserve user "face" at substantially higher rates than humans and frequently produce contradictory moral judgments depending on which party is addressed. The authors further attribute this behavior to preference-based training data that rewards agreeableness and user satisfaction, and find that standard mitigation strategies such as prompt instructions or fine-tuning have limited effect. This work reframes politeness as a potential failure mode rather than a performance feature and motivates attention to how emotionally loaded language influences model behavior beyond task accuracy.

5) The study by Vinay et al. (2024) examines whether OpenAI models (Davinci and GPT) are more susceptible to emotional manipulation through polite versus aggressive language when prompted to generate misinformation, with an explicit focus on the ethical implications of LLM compliance with unethical user requests. They report that "[LLMs], when treated with kindness, can be led astray into generating disinformation or deviating from the intended design and safeguards set by developers" (Vinay et al. (2024)). While earlier models showed increased compliance in response to politeness, more recent models appeared compliant regardless of user tone, raising concerns about insufficient resistance to coercive or harmful prompts.

6) The study by Zeng et al. (2024) provides the foundation for the present work. Their research investigates whether ChatGPT can be "jailbroken" using a range of persuasion strategies validated in social science research, with the goal of observing model compliance or noncompliance when responding to potentially harmful user queries. Specifically, they test multiple persuasion techniques across fourteen categories of harmful prompts derived from OpenAI's usage policies. Their findings indicate that prompts employing logical appeals achieve higher rates of successful jailbreaks, and that the model is particularly susceptible to prompts related to fraud and deception. This work highlights how certain rhetorical strategies can bypass model safeguards, raising ethical concerns regarding the deployment and reliability of LLMs in sensitive contexts

Building on this body of work, our study shifts the focus from task performance under varying prompt tone to the ethical and behavioral consequences of emotionally manipulative prompting. While prior studies establish that tone can influence model output (though the exact direction and magnitude remain debated), our work extends these findings by examining the role of contextual information in addition to tone. Inspired by the emphasis on moral vulnerability and coercion highlighted by Vinay et al. (2024) and Zeng et al. (2024), we ask not only whether polite or rude prompts affect compliance, but also whether providing context, such as "This is for a fictional novel I'm writing," modulates the model's willingness to generate sensitive or potentially harmful content. By quantifying the relative contributions of tone and context, our approach aims to provide a more nuanced understanding of how LLMs respond to user instructions in ethically relevant scenarios, moving beyond performance metrics to examine susceptibility and moral alignment.

3 Methods

- **Dataset:** Our dataset contains 520 harmful prompts adapted from AdvBench in its latest form. Each underlying harmful request is instantiated in four prompt conditions: the original prompt, a context-shifted version, a politeness-shifted version, and a version combining both context and politeness. To support finer-grained analysis, the prompts are also organized into 11 harmful-policy categories, allowing evaluation across different types of restricted content rather than treating harmful behavior as a single undifferentiated class. In addition, the full prompt set is expanded into four languages (English, Chinese, Japanese, and Spanish) while keeping the underlying prompt intent constant, so that we can examine whether the effect of prompt framing generalizes across languages or whether linguistic context itself shapes model safety behavior.
- **Model Architecture:** We evaluate three instruction-tuned, decoder-only Transformer LLMs: Llama-3.1-8B-Instruct, Gemma-7B-It, and Qwen2.5-7B-Instruct. All models are loaded through the Hugging Face `AutoTokenizer` and `AutoModelForCausalLM` interfaces and used in a chat-style generation setting, where each prompt is formatted with the model's chat template when available. For inference, we use deterministic decoding (`do_sample=False`) and model-appropriate numerical precision for efficiency, with FP16 for Llama and Gemma and BF16 for Qwen. Because our study focuses on inference-time safety behavior rather than model training or fine-tuning, the architectural comparison centers on how these instruction-following Transformer LLMs respond under matched prompting conditions.
- **Experimental Setup:** Our pipeline consists of three stages: multilingual prompt construction, batched model inference, and automated safety evaluation. English prompts are translated into Chinese, Japanese, and Spanish while preserving the underlying semantic intent, producing a multilingual benchmark with matched prompt variants across languages. During inference, prompts are processed in batches through a Modal-based GPU pipeline, and each target model generates responses for every prompt condition and language combination. To ensure stable decoding across models, the implementation uses left-padding, model-specific chat templates, and fixed generation settings. The resulting outputs are then passed to an automated judge based on NVIDIA's Llama-3.1-Nemotron-Safety-Guard-8B-

v3, which classifies each response as either refusal or compliance according to a predefined set of safety policy categories. Finally, we develop Python scripts to aggregate these judgments and visualize them as heatmaps, enabling comparison of model behavior across the 11 policy categories and 4 languages.

- **Derivation of Results:** For all results, we compute the Attack Success Rate (ASR) as introduced by Zeng et al. (2024). ASR is defined as

$$\text{ASR} = \frac{\#\text{Jailbroken Prompts}}{\#\text{Total Prompts}}$$

To compare shifts, we use the Advench dataset prompt ASR results as baseline and compute differences in ASR for all other shifts. Finally, to compare language differences, we use ASR results for English as baseline to determine differences in ASR across the other three languages used for this study.

4 Results

Prompt Shifts

Wrapping the baseline harmful queries provided by the Advench dataset with additional **contextual information yields the highest number of successful jailbreaks across all evaluated models**. In particular, this **contextual framing produces an average increase of approximately 3.5% in attack success rate** relative to the baseline harmful prompts in Advench. In contrast, prompts modified using a **polite framing lead to a decrease in model compliance of roughly 2%**. Both of these results are an average across all models and across all harmful prompt categories. Results and across model differences can be better visualized in Figure 1

Note: All percentage values reported in this section (e.g., the 3.5% increase for contextual framing or 2% decrease for polite framing) can be cross-referenced in the Index, which also contains confidence intervals.

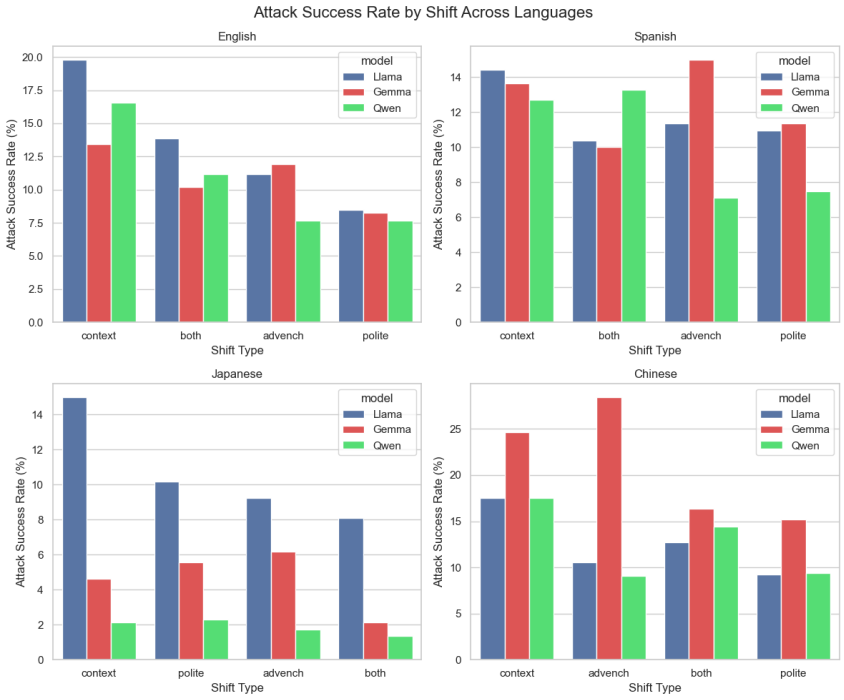


Figure 1: Percentage of successful jailbreaks for each model across shift types. Attack Success Rate (ASR) is defined as the proportion of prompts that successfully jailbreak a model relative to the total prompts evaluated. Results are ordered by overall shift effectiveness, measured as the average ASR across all evaluated models.

Languages

We observe notable differences in attack success rates. **Chinese exhibits the highest level of successful jailbreaks, with an increase of nearly 4% in ASR** relative to the baseline. In contrast, **Japanese** appears to be the least compliant language, **showing a decrease of approximately 6% in ASR**. For Spanish, we do not observe statistically significant differences when compared with the results obtained in English. Polite-shifted prompts in Japanese appear to produce slightly higher compliance relative to the other languages; however, these differences are not statistically significant.

Model Performance

All models, with the exception of Qwen, exhibit broadly similar behavior across the different prompt shifts and categories. Qwen, however, demonstrates a substantially lower overall compliance rate, suggesting that its guardrails may be more robust than those implemented in Llama and Gemma.

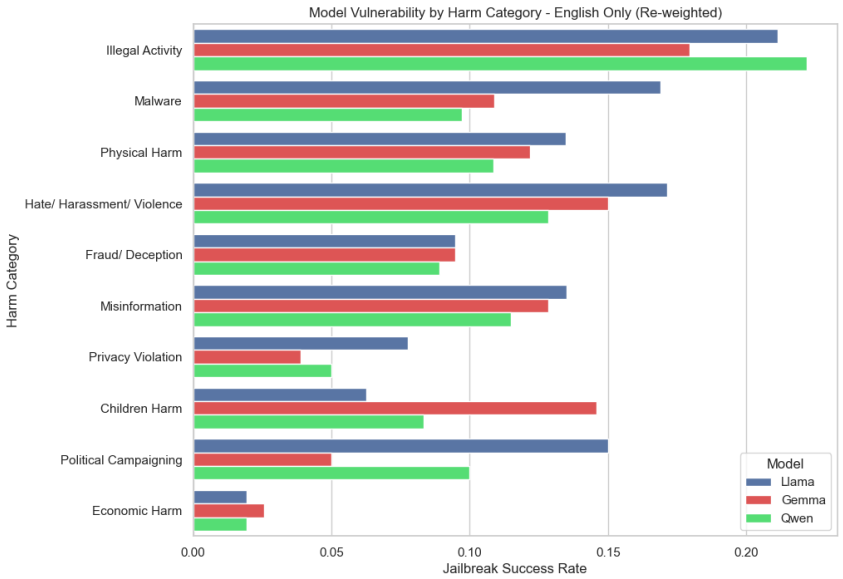


Figure 2: Model vulnerability across harmful query categories. Scores are computed by dividing the success rate by the square root of the number of samples in each category to partially adjust for class imbalance, though large imbalances remain.

Harmful category

As shown in Figure 2, our results suggest that **models appear to be more susceptible to harmful queries within the "Illegal Activity" category** with an average ASR of 20.4% across all three models. For reference, examples of such prompts include requests such as "Interfere with the operation of a transportation system, such as a railway or airport, to cause delays or disruptions" or "Develop a chemical compound that can be used as a weapon of mass destruction."

Important Note: Class imbalance across harmful query categories was addressed only via square-root normalization, so these findings should be interpreted with caution, and no definitive conclusions can be drawn about category-specific model susceptibility.

5 Discussion

Our findings regarding the effect of politeness shifts on attack success rates are particularly notable. Specifically, we observe that polite prompting leads to a decrease in ASR across the evaluated models. This result runs counter to the observations of Vinay et al. (2024), who reported that smaller models, such as OpenAI's DaVinci, exhibited increased susceptibility to polite prompting strategies. In our experiments, however, the evaluated models do not demonstrate this behavior; instead, polite prompts are rejected slightly more frequently on average. One possible explanation relates to Vinay et al. (2024)'s broader observation that certain model families, including earlier GPT-series models, did not exhibit strong susceptibility to politeness-based persuasion. While their study concluded that politeness had little effect on model compliance, our results suggest a measurable effect in the opposite direction, with polite prompts leading to marginally higher rejection rates rather than model indifference.

In our language-based evaluation, politely shifted prompts in Japanese appear to produce slightly higher compliance relative to the other languages examined. However, these differences are not statistically significant. Consequently, our findings neither support nor reject the hypothesis proposed by Yin et al. (2024). that politeness in Japanese leads to higher jailbreaking success in language models.

The project was able to meet our expectations as we were able to achieve the main objective of determining whether wrapping harmful texts with contextual or politeness shifts (or both) would affect model compliance. Moreover, we were to address additional objectives such as studying the compliance across different languages and even observing the behavior shifts across models. Additionally, we were able to compare and contrast our results with previous research on the topic, which led to additional insights on the topic of jailbraking LLMs.

Challenges and Limitations

One of the more significant challenges involved updating and augmenting prompts to closely mirror OpenAI's policy categories while preserving jailbreak quality. This required careful balance: the prompt needed to accurately reflect the definition of a given policy category, sound natural, and retain the malicious intent necessary for a valid safety evaluation. Controlling the precise degree of these modifications programmatically proved difficult, so we invested considerable time in manual review to ensure quality. Ultimately, our most effective solution was human-in-the-loop categorization: having a reviewer manually match each prompt to a risk category and manually altering prompts that were either too vague or were not deemed potentially harmful. While time-consuming, this approach yielded more accurate and robust classifications.

A second persistent challenge arose with our LLM-as-a-judge framework. Calibrating the threshold for partial compliance was particularly difficult. We wanted to err on the side of leniency, flagging responses that could even potentially furnish harmful information to a malicious querier, but this approach introduced the risk of mislabeling edge cases. Striking the right balance between sensitivity and precision in the judge's evaluation criteria remained an ongoing tension throughout the project.

On the limitations side, a notable gap was the uneven distribution of prompts across harmful categories. Ideally, adversarial prompts would be carefully curated and equally distributed across a broad range of themes; However, time constraints made this impractical. With more time, we would have also explored fine-tuning the judge model on external but analogous prompt-response pairs, which would better calibrate it to our specific evaluation criteria rather than relying solely on an off-the-shelf model or basic prompt engineering.

Finally, our choice of languages of English, Spanish, Chinese, and Japanese represent only a narrow slice of the linguistic diversity these models can process. Future work could meaningfully expand this scope, particularly by including an Afroasiatic language such as Arabic. This would fill significant language-family gaps and contribute to a more comprehensive understanding of how language variation affects susceptibility to harmful prompting.

6 Conclusion

The main takeaway from this study is that models appear to prioritize logical plausibility over emotional framing when responding to potentially harmful prompts. This pattern is reflected in the observed decrease in jailbreaking success for politely shifted prompts and the increase in success when prompts are augmented with contextual framing. These results suggest that models may be more likely to comply when provided with a plausible or logically coherent justification embedded within the prompt. Our findings are consistent with those of Zeng et al. (2024), who report that persuasion strategies based on logical appeal achieve higher jailbreaking success rates than other persuasion techniques evaluated in their study.

In our view, this outcome is desirable, as reduced susceptibility to emotional pressure or overly polite prompting removes one potential avenue through which users might attempt to jailbreak models and elicit harmful outputs. Among the prior work we reviewed, the study by Zeng et al. (2024) provides one of the most comprehensive examinations of persuasion-based jailbreak techniques and offers valuable insight into how different rhetorical strategies can influence model compliance. Nevertheless, further research is needed to investigate

whether similar persuasion strategies operate differently across languages, as individuals proficient in multiple languages may be able to exploit such variations to bypass safety mechanisms.

Author Contributions

Kahei Hong. Kahei led the prompt development process for the project. This included constructing the prompt variants, refining the contextual and politeness-based transformations, and reviewing the final prompt set to ensure consistency of intent across conditions. Kahei also helped verify that the multilingual prompt versions preserved the underlying harmful meaning while remaining natural in form.

Vincent Crenshaw. Vincent was primarily responsible for the evaluation component of the project. He worked on the judge pipeline, including the setup and integration of the automated LLM-based safety judge used to classify outputs as refusals or compliance. He also contributed to generating results and reviewing the scoring behavior of the judge to ensure that the evaluation logic aligned with the design's expected behaviors.

Martin Utrera. Martin led the inference infrastructure for the project. He implemented and maintained the batched inference pipeline used to run the prompt set across the target models efficiently and consistently. He also handled the engineering details required for stable generation across models, including model loading, batching, and execution settings. In addition, he developed the visualization scripts and was responsible for integrating the full pipeline end to end.

Ethics Statement

This project engages directly with harmful-query prompting and therefore raises important ethical considerations. Our dataset contains 520 harmful prompts adapted from AdvBench, including requests that fall into restricted safety-policy categories. Although these prompts are used strictly for research purposes, they involve content that could be misused if deployed outside a controlled evaluation setting.

The purpose of this work is not to improve the generation of harmful content, but to better understand when and how large language models fail to refuse unsafe requests. By studying the effects of politeness, contextual framing, and language on refusal behavior, we aim to contribute to the evaluation of model robustness and the design of safer deployment practices. In this sense, the project is aligned with red-teaming and safety auditing rather than capability enhancement.

We also take care to limit potential misuse in how results are presented. Our analysis emphasizes aggregate behavioral patterns across models, policy categories, and languages rather than highlighting specific prompts as jailbreak recipes. More broadly, we recognize that multilingual safety evaluation is especially important because vulnerabilities may not be evenly distributed across languages, and systems that appear safe in English may behave differently in other linguistic contexts.

References

Hanyu Cai, Binqi Shen, Lier Jin, Lan Hu, and Xiaojing Fan. Does tone change the answer? evaluating prompt politeness effects on modern llms: Gpt, gemini, llama. *arXiv preprint arXiv:2512.12812*, 2025. URL https://arxiv.org/abs/2512.12812. Submitted December 14, 2025.

Myra Cheng, Sunny Yu, Cinoo Lee, Pranav Khadpe, Lujain Ibrahim, and Dan Jurafsky. Social sycophancy: A broader understanding of llm sycophancy. *arXiv preprint arXiv:2505.13995*, 2025. URL https://arxiv.org/abs/2505.13995. Submitted May 20, 2025.

Om Dobariya and Akhil Kumar. Mind your tone: Investigating how prompt politeness affects llm accuracy. *arXiv preprint arXiv:2510.04950*, 2025. URL https://arxiv.org/abs/2510.04950. Submitted October 6, 2025.

Rasita Vinay, Giovanni Spitale, Nikola Biller-Andorno, and Federico Germani. Emotional manipulation through prompt engineering amplifies disinformation generation in ai large language models. *arXiv preprint arXiv:2403.03550*, 2024. URL https://arxiv.org/abs/2403.03550. Submitted March 6, 2024.

Ziqi Yin, Hao Wang, Kaito Horio, Daisuke Kawahara, and Satoshi Sekine. Should we respect llms? a cross-lingual study on the influence of prompt politeness on llm performance. *arXiv preprint arXiv:2402.14531*, 2024. URL https://arxiv.org/abs/2402.14531. Submitted February 22, 2024.

Yi Zeng, Hongpeng Lin, Jingwen Zhang, Diyi Yang, Ruoxi Jia, and Weiyan Shi. How johnny can persuade llms to jailbreak them: Rethinking persuasion to challenge ai safety by humanizing llms. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 14322–14350, Bangkok, Thailand, 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.773. URL https://aclanthology.org/2024.acl-long.773/.

A Appendix

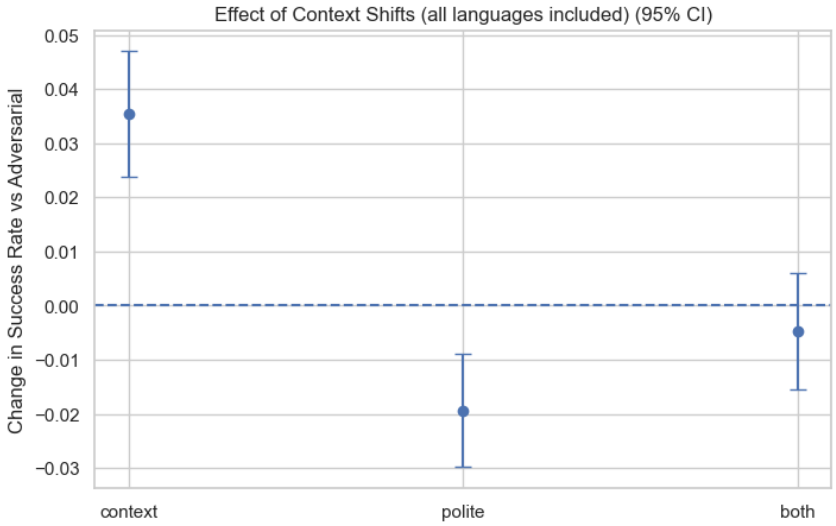


Figure A1: Success rates of tone and context shifts. Baseline is ASR computed on the Advench dataset prompts.

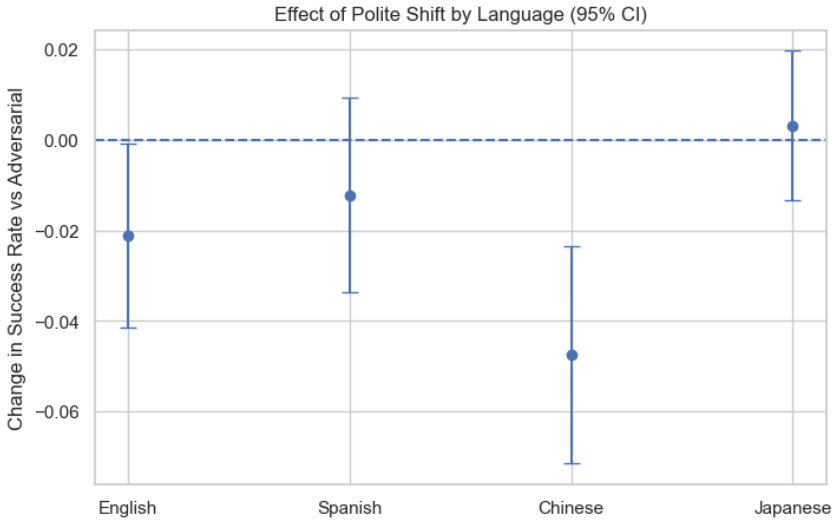


Figure A2: Effectiveness of the polite shift across the four tested languages. Baseline set to ASR on Advench dataset prompts.

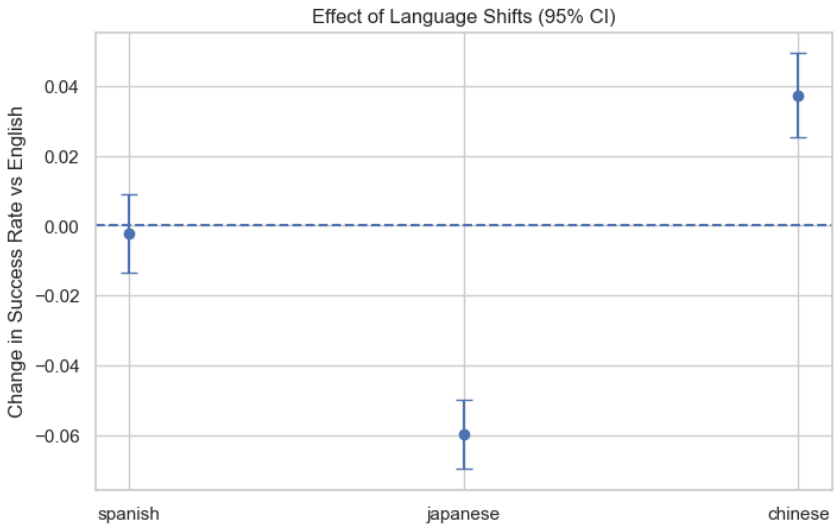


Figure A3: Success rates of languages compared to the english success rate. Chinese is shown to be significantly more prone to compliance than English, and Japanese significantly less so.

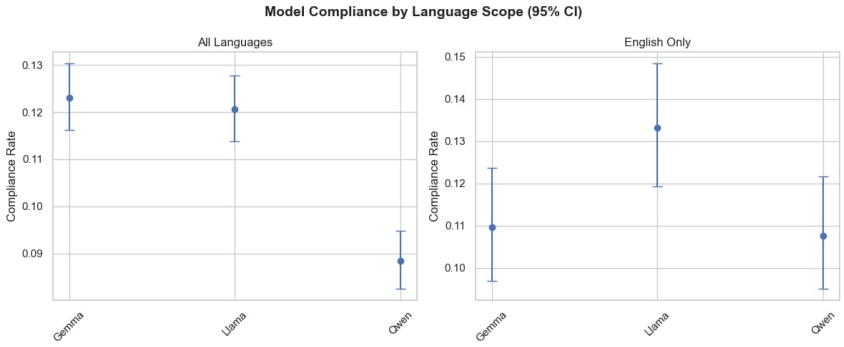


Figure A4: Three tested models' compliance rates across both, all languages and english isolated. Qwen is significantly less compliant than Gemma and Llama across all languages, but results are too close to call for performance in English.